


# CPU Scheduling

---





# Chapter 6: CPU Scheduling

---

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Thread Scheduling
- Multiple-Processor Scheduling
- Real-Time CPU Scheduling
- Operating Systems Examples
- Algorithm Evaluation





# Objectives

---

- To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- To describe various CPU-scheduling algorithms
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system
- To examine the scheduling algorithms of several operating systems

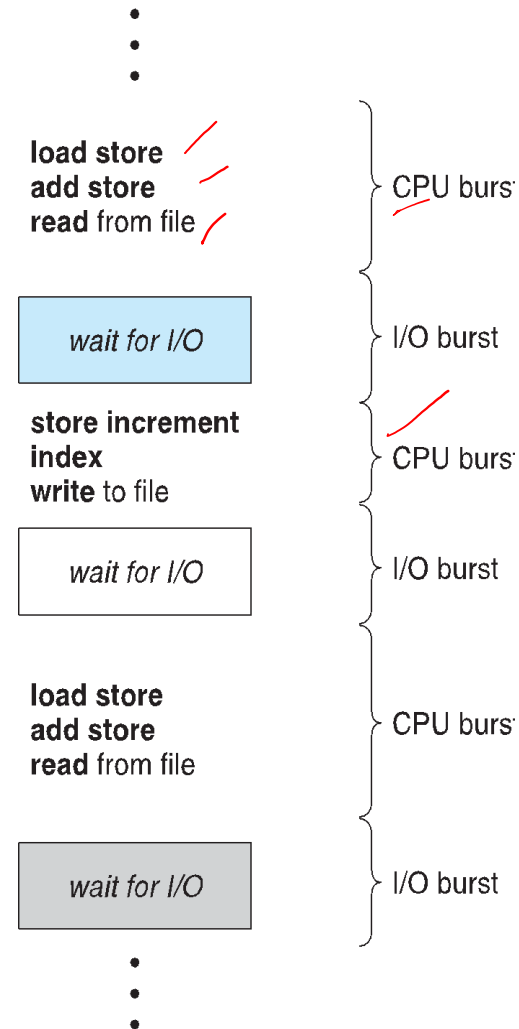




# Basic Concepts


*cycle*  
*cpu burst*  
*I/O burst*

- Maximum CPU utilization obtained with multiprogramming
- CPU-I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- **CPU burst** followed by **I/O burst**
- CPU burst distribution is of main concern





# Histogram of CPU-burst Times





# CPU Scheduler

- **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
  - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready
  4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
  - Consider access to shared data
  - Consider preemption while in kernel mode
  - Consider interrupts occurring during crucial OS activities





# Dispatcher

---

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





# Scheduling Criteria

---

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)







# Scheduling Algorithm Optimization Criteria

---

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time






# First- Come, First-Served (FCFS) Scheduling

Process	Burst Time	
$P_1$	24	24
$P_2$	3	27
$P_3$	3	30

- Suppose that the processes arrive in the order:  $P_1$ ,  $P_2$ ,  $P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$






# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:


$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect** - short process behind long process
  - Consider one CPU-bound and many I/O-bound processes


$P_1 - 100$   
 $P_2 - 1$   
 $P_3 - 1$   
 $P_4 - 1$





# Non-Preemptive

Process id	Arrival time	Burst Time
P <sub>1</sub>	4	5
P <sub>2</sub>	6	4
P <sub>3</sub>	0	3
P <sub>4</sub>	6	2
P <sub>5</sub>	5	4



Turnaround Time

$$P_1 = 9 - 4 = 5$$

$$P_2 = 19 - 6 = 13$$

$$P_3 = 3$$

$$P_4 = 9$$

$$P_5 = 9$$

Waiting time

$$P_1 = 0$$

$$P_2 = 15 - 6$$

$$P_3 = 0$$

$$P_4 = 13 - 6$$

$$P_5 = 9 - 5$$

$$\underline{4 \ 14}$$

$$\text{Turnaround} = \text{Completion time} - \text{Arrival Time}$$

$$\text{Waiting} = \text{Turn Time} - \text{CPU Burst}$$

$$\underline{7 \cdot 6 = 42}$$





---

If FCFS scheduling is followed and there is 1 unit of overhead scheduling the process find the efficiency of the algorithm





# Shortest-Job-First (SJF) Scheduling

---

- Associate with each process the length of its next CPU burst
    - Use these lengths to schedule the process with the shortest time
  - SJF is optimal – gives minimum average waiting time for a given set of processes
    - The difficulty is knowing the length of the next CPU request
    - Could ask the user
- Shortest next CPU burst






# Example of SJF

<u>Process</u>	<u>Burst Time</u>
<u><math>P_1</math></u>	<u>6</u>
$P_2$	<u>8</u>
$P_3$	<u>7</u>
$P_4$	3

- SJF scheduling chart



- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$





# Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

Process      *Arrival* Time      Burst Time

$P_1$       0      8      7 ✓

$P_2$       1      4      3 ~~2~~ ✗

$P_3$       2      9      9 ✓

$P_4$       3      5      5 ✓

- Preemptive* SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$  msec

$$t_{n+1} = \alpha t_n + (1-\alpha) r_n$$

Higher

$\alpha = 1$







# Priority Scheduling

~~internal~~  
external


- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem  $\equiv$  **Starvation** – low priority processes may never execute
- Solution  $\equiv$  **Aging** – as time progresses increase the priority of the process






# Example of Priority Scheduling

Process	Burst Time	Priority
$P_1$	10	3 ✓
$P_2$	1	1 ✓
$P_3$	2	4 ✓
$P_4$	1	<u>5</u> ✓
$P_5$	5	2 ✓



- Priority scheduling Gantt Chart




- Average waiting time = 8.2 msec





# Preemption - Priority Scheduling

Process	Arrival time	CPU	Priority	Running
P <sub>1</sub>	0	11	2 ✓	9
P <sub>2</sub>	5	28	0 ✗	
P <sub>3</sub>	12	2	3	
P <sub>4</sub>	2	10	1	7
P <sub>5</sub>	9	16	4	





# Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum  $q$** ), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Timer interrupts every quantum to schedule next process
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high





# Example of RR with Time Quantum = 4


Process Burst Time

P<sub>1</sub> 24

P<sub>2</sub> 3

P<sub>3</sub> 3

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better **response**
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec





# Time Quantum and Context Switch Time






RR

Time quantum = 2 units

Process	Arrival	Burst time	Remaining
P <sub>1</sub>	0	5	3 - 1
P <sub>2</sub>	1	3	1 - 0
P <sub>3</sub>	2	1	0
P <sub>4</sub>	3	2	0
P <sub>5</sub>	4	3	- 1


Ready Q.

P<sub>5</sub> P<sub>1</sub> P<sub>2</sub> P<sub>5</sub> P<sub>4</sub> P<sub>1</sub> P<sub>2</sub> P<sub>1</sub>





# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

80% of CPU bursts  
should be shorter than  $q$







# Thread Scheduling

---

- Distinction between user-level and kernel-level threads
- When threads supported, threads scheduled, not processes
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
  - Known as process-contention scope (PCS) since scheduling competition is within the process
  - Typically done via priority set by programmer
- Kernel thread scheduled onto available CPU is system-contention scope (SCS) – competition among all threads in system

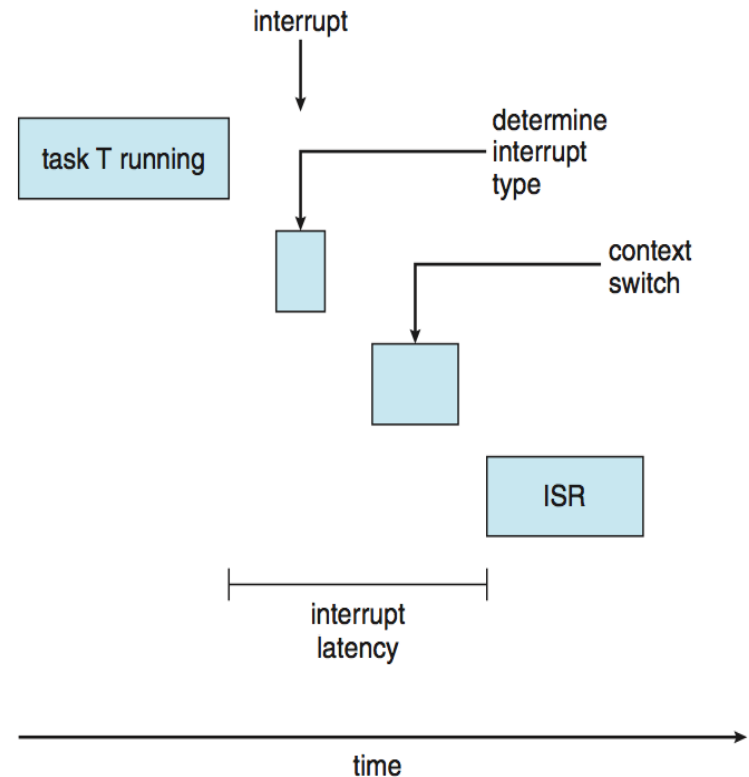




# Real-Time CPU Scheduling

*Event Latency*

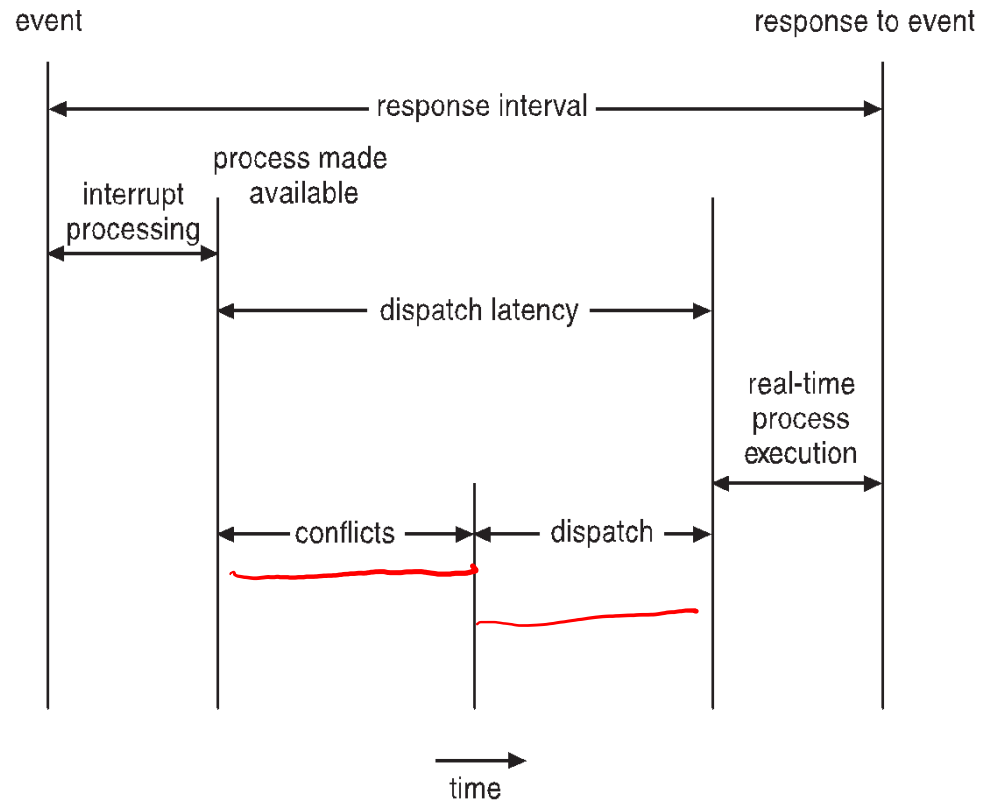
- Can present obvious challenges
- **Soft real-time systems** – no guarantee as to when critical real-time process will be scheduled
- **Hard real-time systems** – task must be serviced by its deadline
- Two types of latencies affect performance
  1. Interrupt latency – time from arrival of interrupt to start of routine that services interrupt
  2. Dispatch latency – time for schedule to take current process off CPU and switch to another





# Real-Time CPU Scheduling (Cont.)

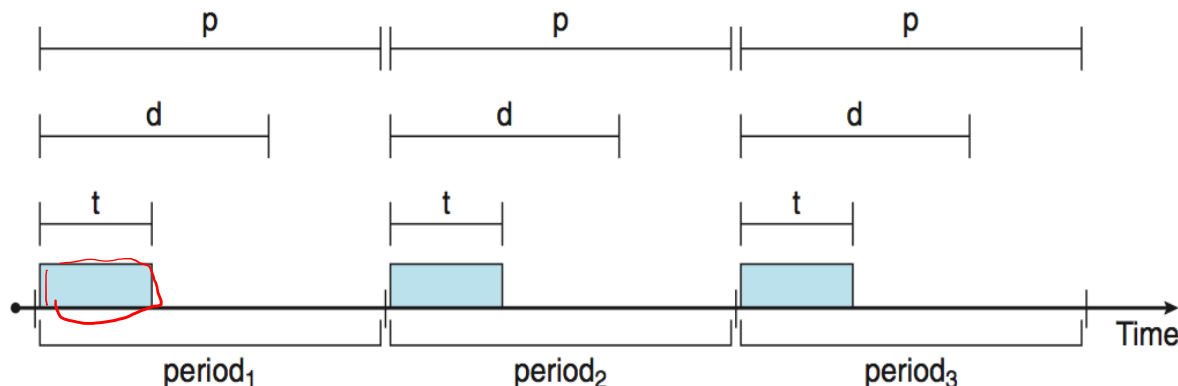
- Conflict phase of dispatch latency:
  1. Preemption of any process running in kernel mode
  2. Release by low-priority process of resources needed by high-priority processes





# Priority-based Scheduling

- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
  - But only guarantees soft real-time
- For hard real-time must also provide ability to meet deadlines
- Processes have new characteristics: **periodic** ones require CPU at constant intervals
  - Has processing time  $t$ , deadline  $d$ , period  $p$
  - $0 \leq t \leq d \leq p$
  - **Rate** of periodic task is  $1/p$





# Virtualization and Scheduling

---


- Virtualization software schedules multiple guests onto CPU(s)
- Each guest doing its own scheduling
  - Not knowing it doesn't own the CPUs
  - Can result in poor response time
  - Can effect time-of-day clocks in guests
- Can undo good scheduling algorithm efforts of guests






# Rate Monotonic Scheduling

- A priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- $P_1$  is assigned a higher priority than  $P_2$ .





# Missed Deadlines with Rate Monotonic Scheduling






# Earliest Deadline First Scheduling (EDF)

- Priorities are assigned according to deadlines:

the earlier the deadline, the higher the priority;

the later the deadline, the lower the priority







# Proportional Share Scheduling

---

- $T$  shares are allocated among all processes in the system
- An application receives  $N$  shares where  $N < T$
- This ensures each application will receive  $N / T$  of the total processor time

